

COMP111: Artificial Intelligence

Section 6. Uniform cost search and informed (or heuristic) tree search

Frank Wolter

Recap

- ▶ Basic problem solving techniques:
 - ▶ **Breadth-first search**
complete but expensive.
 - ▶ **Depth-first search**
cheap but incomplete
- ▶ Variations and combinations:
 - ▶ **Limited depth search**
 - ▶ **Iterative deepening search**
 - ▶ **Avoiding repeated states**
 - ▶ **Bi-directional search**

Overview

- ▶ Introduce **uniform cost search**: generalizing breadth-first search to search problems with costs.

Overview

- ▶ Introduce **uniform cost search**: generalizing breadth-first search to search problems with costs.
- ▶ Introduce **heuristics**: rules of thumb

Overview

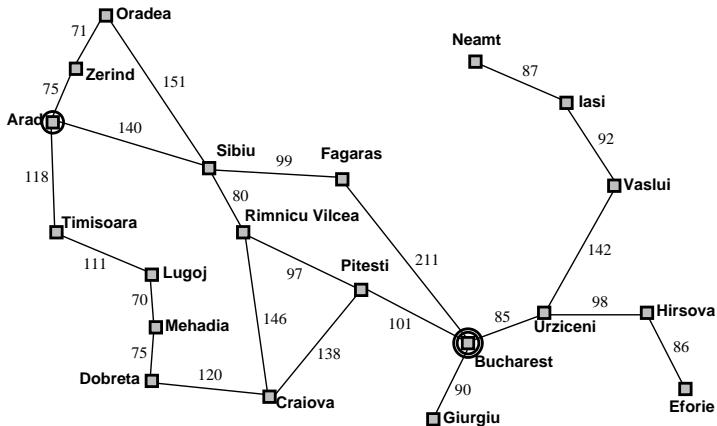
- ▶ Introduce **uniform cost search**: generalizing breadth-first search to search problems with costs.
- ▶ Introduce **heuristics**: rules of thumb
- ▶ Introduce **heuristic search**:
 - ▶ greedy search
 - ▶ A* search

Search graph with costs

- ▶ A **path cost function**,

g : Paths \rightarrow real numbers

gives a **cost** to each path. We assume that the cost of a path is the sum over the costs of the steps in the path.



Uniform Cost Search

- ▶ Why not expand the **cheapest** path first?

Uniform Cost Search

- ▶ Why not expand the **cheapest** path first?
- ▶ Intuition: cheapest is likely to be best!

Uniform Cost Search

- ▶ Why not expand the **cheapest** path first?
- ▶ Intuition: cheapest is likely to be best!
- ▶ Performance is like breadth-first search but we select (expand) the minimum cost path rather than the shortest path.

Uniform Cost Search

- ▶ Why not expand the **cheapest** path first?
- ▶ Intuition: cheapest is likely to be best!
- ▶ Performance is like breadth-first search but we select (expand) the minimum cost path rather than the shortest path.
- ▶ Uniform cost search behaves in exactly the same way as breadth-first search if the cost of every step is the same.


General Algorithm for Uniform Cost Search

- 1: **Input:** a start state s_0
- 2: for each state s the successors of s
- 3: a test $\text{goal}(s)$ checking whether s is a goal state
- 4: $g(s_0 \dots s_k)$ for every path $s_0 \dots s_k$
- 5:
- 6: Set $\text{frontier} := \{s_0\}$
- 7: **while** frontier is not empty **do**
- 8: select and remove from frontier the path $s_0 \dots s_k$
- 9: with $g(s_0 \dots s_k)$ minimal
- 10: **if** $\text{goal}(s_k)$ **then**
- 11: return $s_0 \dots s_k$ (and terminate)
- 12: **else** for every successor s of s_k add $s_0 \dots s_k s$ to frontier
- 13: **end if**
- 14: **end while**

Uniform Cost Example

Reaching G from S


Exp. paths	Frontier
	{S : 0}



Uniform Cost Example

Reaching G from S

Exp. paths	Frontier
	$\{S : 0\}$
S not goal	$\{SA : 5, SB : 2, SC : 4\}$



Selected path: $S : 0$


Is the last state in S_{goal} ? No

Expand S: add $SA : 0 + 5$, $SB : 0 + 2$, $SC : 0 + 4$ to the frontier

Uniform Cost Example

Reaching G from S

Exp. paths	Frontier
	$\{S : 0\}$
S not goal	$\{SA : 5, SB : 2, SC : 4\}$
SB not goal	$\{SA : 5, SC : 4, SBG : 8\}$



Selected path: $SB : 2$


Is the last state in S_{goal} ? No

Expand SB : add $SBG : 2 + 6$ to the frontier

Uniform Cost Example

Reaching G from S

Exp. paths	Frontier
	$\{S : 0\}$
S not goal	$\{SA : 5, SB : 2, SC : 4\}$
SB not goal	$\{SA : 5, SC : 4, SBG : 8\}$
SC not goal	$\{SA : 5, SBG : 8, SCF : 6\}$



Selected path: $SC : 4$


Is the last state in S_{goal} ? No

Expand SC : add $SCF : 4 + 2$ to the frontier

Uniform Cost Example

Reaching G from S

Exp. paths	Frontier
	$\{S : 0\}$
S not goal	$\{SA : 5, SB : 2, SC : 4\}$
SB not goal	$\{SA : 5, SC : 4, SBG : 8\}$
SC not goal	$\{SA : 5, SBG : 8, SCF : 6\}$
SA not goal	$\{SBG : 8, SCF : 6, SAD : 14, SAE : 9\}$



Selected path: $SA : 5$


Is the last state in S_{goal} ? No

Expand SA : add $SAD : 5 + 9$ and $SAE : 5 + 4$

Uniform Cost Example

Reaching G from S

Exp. paths	Frontier
	$\{S : 0\}$
S not goal	$\{SA : 5, SB : 2, SC : 4\}$
SB not goal	$\{SA : 5, SC : 4, SBG : 8\}$
SC not goal	$\{SA : 5, SBG : 8, SCF : 6\}$
SA not goal	$\{SBG : 8, SCF : 6, SAD : 14, SAE : 9\}$
SCF not goal	$\{SBG : 8, SAD : 14, SAE : 9, SCFG : 7\}$



Selected path: $SCF : 6$


Is the last state in S_{goal} ? No

Expand SCF : add $SCFG : 6 + 1$ to the frontier

Uniform Cost Example

Reaching G from S

Exp. paths	Frontier
	$\{S : 0\}$
S not goal	$\{SA : 5, SB : 2, SC : 4\}$
SB not goal	$\{SA : 5, SC : 4, SBG : 8\}$
SC not goal	$\{SA : 5, SBG : 8, SCF : 6\}$
SA not goal	$\{SBG : 8, SCF : 6, SAD : 14, SAE : 9\}$
SCF not goal	$\{SBG : 8, SAD : 14, SAE : 9, SCFG : 7\}$
$SCFG$ goal	$\{SBG : 8, SAD : 14, SAE : 9\}$



Selected path: $SCFG : 7$

Is the last state in S_{goal} ? Yes!

Path found: $SCFG$ with a cost of 7

Properties of Uniform Cost Search

- ▶ Complete and optimal: Uniform cost search guaranteed to find cheapest solution **assuming path costs grow monotonically**, i.e. the cost of a path increases if we move along it.
- ▶ In other words, we assume that adding another step to a path makes it more costly, i.e. $g(s_0 \dots s_k) < g(s_0 \dots s_k s)$.
- ▶ If path costs **don't** grow monotonically, then exhaustive search is required.
- ▶ Time and space complexity: the same as breadth first search.

Proof of Optimality of Uniform Cost

Theorem. Uniform Cost is optimal if the cost of each action is positive and the cost of a path is the sum of the costs of its actions.

Proof of Optimality of Uniform Cost

Theorem. Uniform Cost is optimal if the cost of each action is positive and the cost of a path is the sum of the costs of its actions.

Proof. Let $s_0 \cdots s_k$ be a cheapest path to a goal state s_k . We prove that the following always holds:

- (A) When entering the while loop, the frontier contains an initial part $s_0 \cdots s_i$ of $s_0 \cdots s_k$ (in other words, $i \leq k$). (In particular that the cost of the selected path does not exceed the cost of $s_0 \cdots s_k$ (so its cost is $\leq g(s_0 \cdots s_k)$)).

Proof of Optimality of Uniform Cost

Theorem. Uniform Cost is optimal if the cost of each action is positive and the cost of a path is the sum of the costs of its actions.

Proof. Let $s_0 \cdots s_k$ be a cheapest path to a goal state s_k . We prove that the following always holds:

(A) When entering the while loop, the frontier contains an initial part $s_0 \cdots s_i$ of $s_0 \cdots s_k$ (in other words, $i \leq k$). (In particular that the cost of the selected path does not exceed the cost of $s_0 \cdots s_k$ (so its cost is $\leq g(s_0 \cdots s_k)$)).

Proof by Induction over the number of times the while loop is entered.

Proof of Optimality of Uniform Cost

Theorem. Uniform Cost is optimal if the cost of each action is positive and the cost of a path is the sum of the costs of its actions.

Proof. Let $s_0 \cdots s_k$ be a cheapest path to a goal state s_k . We prove that the following always holds:

(A) When entering the while loop, the frontier contains an initial part $s_0 \cdots s_i$ of $s_0 \cdots s_k$ (in other words, $i \leq k$). (In particular that the cost of the selected path does not exceed the cost of $s_0 \cdots s_k$ (so its cost is $\leq g(s_0 \cdots s_k)$)).

Proof by Induction over the number of times the while loop is entered.

(A) is true in the beginning when the frontier only contains the start state s_0 .

Proof of Optimality of Uniform Cost

Theorem. Uniform Cost is optimal if the cost of each action is positive and the cost of a path is the sum of the costs of its actions.

Proof. Let $s_0 \cdots s_k$ be a cheapest path to a goal state s_k . We prove that the following always holds:

(A) When entering the while loop, the frontier contains an initial part $s_0 \cdots s_i$ of $s_0 \cdots s_k$ (in other words, $i \leq k$). (In particular that the cost of the selected path does not exceed the cost of $s_0 \cdots s_k$ (so its cost is $\leq g(s_0 \cdots s_k)$)).

Proof by Induction over the number of times the while loop is entered.

(A) is true in the beginning when the frontier only contains the start state s_0 .

Now assume for a proof by induction that (A) is true for a frontier F we have reached and we enter the while loop again.

Proof of Optimality of Uniform Cost

By induction hypothesis, the frontier F contains an initial part $s_0 \cdots s_j$ of $s_0 \cdots s_k$. We show that in the next run through the while loop either a path of cost $\leq g(s_0 \cdots s_k)$ is returned or (A) holds for the updated frontier.

Proof of Optimality of Uniform Cost

By induction hypothesis, the frontier F contains an initial part $s_0 \cdots s_i$ of $s_0 \cdots s_k$. We show that in the next run through the while loop either a path of cost $\leq g(s_0 \cdots s_k)$ is returned or (A) holds for the updated frontier.

If $i < k$ and so $s_0 \cdots s_i$ is shorter than $s_0 \cdots s_k$, then (A) still holds for the updated frontier after the while loop:

First, if $s_0 \cdots s_i$ itself is selected then (A) holds as $s_0 \cdots s_{i+1}$ is added to the frontier.

If another path is selected then $s_0 \cdots s_i$ is still in the updated frontier and so (A) holds.

Proof of Optimality of Uniform Cost

By induction hypothesis, the frontier F contains an initial part $s_0 \cdots s_i$ of $s_0 \cdots s_k$. We show that in the next run through the while loop either a path of cost $\leq g(s_0 \cdots s_k)$ is returned or (A) holds for the updated frontier.

If $i < k$ and so $s_0 \cdots s_i$ is shorter than $s_0 \cdots s_k$, then (A) still holds for the updated frontier after the while loop:

First, if $s_0 \cdots s_i$ itself is selected then (A) holds as $s_0 \cdots s_{i+1}$ is added to the frontier.

If another path is selected then $s_0 \cdots s_i$ is still in the updated frontier and so (A) holds.

On the other hand, if $i = k$, then we are done if $s_0 \cdots s_k$ is selected as $s_0 \cdots s_k$ is the output then.

If another path is selected then its cost is $\leq g(s_0 \cdots s_k)$.

So if it is a path to the goal it is not more expensive than $s_0 \cdots s_k$ and the output is optimal.

If it is not a path to the goal then (A) still holds as $s_0 \cdots s_k$ is in the updated frontier.

Real Life Problems

- ▶ Whatever search technique we use, **exponential time complexity**.

Real Life Problems

- ▶ Whatever search technique we use, **exponential time complexity**.
- ▶ Tweaks to the algorithm will not reduce this to polynomial.

Real Life Problems

- ▶ Whatever search technique we use, **exponential time complexity**.
- ▶ Tweaks to the algorithm will not reduce this to polynomial.
- ▶ We need problem specific knowledge to **guide** the search.

Real Life Problems

- ▶ Whatever search technique we use, **exponential time complexity**.
- ▶ Tweaks to the algorithm will not reduce this to polynomial.
- ▶ We need problem specific knowledge to **guide** the search.
- ▶ Simplest form of problem specific knowledge is **heuristic**.

Real Life Problems

- ▶ Whatever search technique we use, **exponential time complexity**.
- ▶ Tweaks to the algorithm will not reduce this to polynomial.
- ▶ We need problem specific knowledge to **guide** the search.
- ▶ Simplest form of problem specific knowledge is **heuristic**.
- ▶ Standard implementation in search is via an **evaluation function** which indicates desirability of selecting (expanding) state.

Informed Strategies

- ▶ Use problem-specific knowledge to make the search more efficient.
- ▶ Idea: based on your knowledge, select the most promising path first.
- ▶ Rather than trying all possible search paths, you try to focus on paths that get you nearer to the goal state according to your estimate.

Heuristics

- ▶ Consider heuristics that estimate the cost of cheapest path from a state to a goal state.

Heuristics

- ▶ Consider heuristics that estimate the cost of cheapest path from a state to a goal state.
- ▶ We have a **heuristic function**,

$$h : \text{States} \rightarrow \text{real numbers}$$

which estimates the cost of going from that state to the goal.
 h can be any function but $h(s) = 0$ if s is a goal.

Heuristics

- ▶ Consider heuristics that estimate the cost of cheapest path from a state to a goal state.
- ▶ We have a **heuristic function**,

$$h : \text{States} \rightarrow \text{real numbers}$$

which estimates the cost of going from that state to the goal.
 h can be any function but $h(s) = 0$ if s is a goal.

- ▶ Example: In route finding, heuristic might be straight line distance from node to destination.

Heuristics

- ▶ Consider heuristics that estimate the cost of cheapest path from a state to a goal state.
- ▶ We have a **heuristic function**,

$$h : \text{States} \rightarrow \text{real numbers}$$

which estimates the cost of going from that state to the goal.
 h can be any function but $h(s) = 0$ if s is a goal.

- ▶ Example: In route finding, heuristic might be straight line distance from node to destination.
- ▶ Greedy search: expands the path that **appears** to be closest to goal.


General algorithm for greedy search

- 1: **Input:** a start state s_0
- 2: for each state s the successors of s
- 3: a test $\text{goal}(s)$ checking whether s is a goal state
- 4: $h(s)$ for every state s
- 5:
- 6: Set $\text{frontier} := \{s_0\}$
- 7: **while** frontier is not empty **do**
- 8: select and remove from frontier the path $s_0 \dots s_k$
- 9: with $h(s_k)$ minimal
- 10: **if** $\text{goal}(s_k)$ **then**
- 11: return $s_0 \dots s_k$ (and terminate)
- 12: **else** for every successor s of s_k add $s_0 \dots s_k s$ to frontier
- 13: **end if**
- 14: **end while**

Greedy Example

Reaching G from S


Exp. paths	Frontier
	$\{S : 8\}$



Greedy Example

Reaching G from S

Exp. paths	Frontier
	$\{S : 8\}$
S not goal	$\{SA : 8, SB : 4, SC : 3\}$



Selected path: $S : 8$


Is the last state in S_{goal} ? No

Expand S : add $SA : 8, SB : 4$, and $SC : 3$ to the frontier

Greedy Example

Reaching G from S

Exp. paths	Frontier
	$\{S : 8\}$
S not goal	$\{SA : 8, SB : 4, SC : 3\}$
SC not goal	$\{SA : 8, SB : 4, SCG : 0\}$



Selected path: $SC : 3$


Is the last state in S_{goal} ? No

Expand SC : add $SCG : 0$ to the frontier

Greedy Example

Reaching G from S

Exp. paths	Frontier
	$\{S : 8\}$
S not goal	$\{SA : 8, SB : 4, SC : 3\}$
SC not goal	$\{SA : 8, SB : 4, SCG : 0\}$
SCG goal	$\{SA : 8, SB : 4\}$




Selected path: $SCG : 0$

Is the last state in S_{goal} ? Yes!

Path found: SCG with a cost of 13


Romania Example




Straight-line distance
to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374


Greedy Search Example



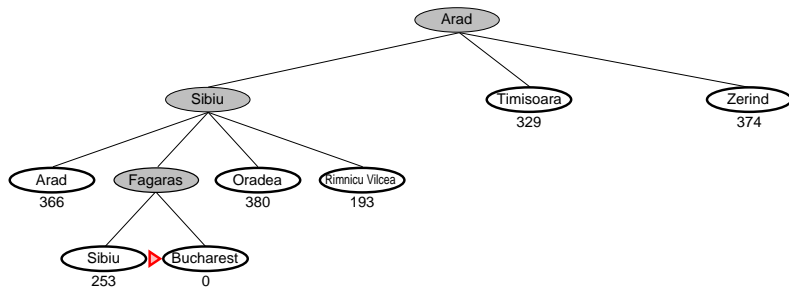
Greedy Search Example



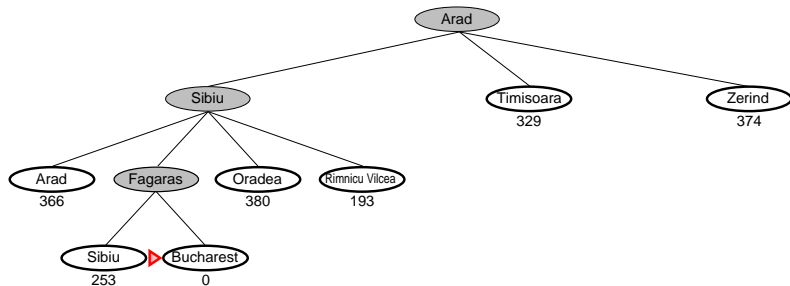
Greedy Search Example



Greedy Search Example



Greedy Search Example



Total distance to go: 450 km

Properties of Greedy Search

- ▶ Greedy search sometimes finds solutions quickly.

Properties of Greedy Search

- ▶ Greedy search sometimes finds solutions quickly.
- ▶ Doesn't always find best.

Properties of Greedy Search

- ▶ Greedy search sometimes finds solutions quickly.
- ▶ Doesn't always find best.
- ▶ May not find a solution if there is one (incomplete).

Properties of Greedy Search

- ▶ Greedy search sometimes finds solutions quickly.
- ▶ Doesn't always find best.
- ▶ May not find a solution if there is one (incomplete).
- ▶ Susceptible to false starts.

Properties of Greedy Search

- ▶ Greedy search sometimes finds solutions quickly.
- ▶ Doesn't always find best.
- ▶ May not find a solution if there is one (incomplete).
- ▶ Susceptible to false starts.
- ▶ Only looking at current state. Ignores past!

A* Search

- ▶ A* was developed around 1968 in Stanford by the team that constructed Shakey, the robot. You might want to watch the video on canvas about Shakey.

A* Search

- ▶ A* was developed around 1968 in Stanford by the team that constructed Shakey, the robot. You might want to watch the video on canvas about Shakey.
- ▶ Basic idea is to **combine** uniform cost search **and** greedy search.

A* Search

- ▶ A* was developed around 1968 in Stanford by the team that constructed Shakey, the robot. You might want to watch the video on canvas about Shakey.
- ▶ Basic idea is to **combine** uniform cost search **and** greedy search.
- ▶ We look at the **cost so far** and the **estimated cost to goal**.

A* Search

- ▶ A* was developed around 1968 in Stanford by the team that constructed Shakey, the robot. You might want to watch the video on canvas about Shakey.
- ▶ Basic idea is to **combine** uniform cost search **and** greedy search.
- ▶ We look at the **cost so far** and the **estimated cost to goal**.
- ▶ Thus, we use **heuristic** f :

$$f(s_0 \dots s_k) = g(s_0 \dots s_k) + h(s_k)$$

where

- ▶ $g(s_0 \dots s_k)$ is path cost of $s_0 \dots s_k$;
- ▶ $h(s_k)$ is expected cost of cheapest solution from s_k .

A* Search

- ▶ A* was developed around 1968 in Stanford by the team that constructed Shakey, the robot. You might want to watch the video on canvas about Shakey.
- ▶ Basic idea is to **combine** uniform cost search **and** greedy search.
- ▶ We look at the **cost so far** and the **estimated cost to goal**.
- ▶ Thus, we use **heuristic** f :

$$f(s_0 \dots s_k) = g(s_0 \dots s_k) + h(s_k)$$

where

- ▶ $g(s_0 \dots s_k)$ is path cost of $s_0 \dots s_k$;
 - ▶ $h(s_k)$ is expected cost of cheapest solution from s_k .
- ▶ Aims to minimise **overall cost**.

General algorithm for A* search


- 1: **Input:** a start state s_0
- 2: for each state s the successors of s
- 3: a test $\text{goal}(s)$ checking whether s is a goal state
- 4: $g(s_0 \dots s_k)$ for every path $s_0 \dots s_k$
- 5: $h(s)$ for every state s
- 6:
- 7: Set $\text{frontier} := \{s_0\}$
- 8: **while** frontier is not empty **do**
- 9: select and remove from frontier the path $s_0 \dots s_k$
- 10: with $g(s_0 \dots s_k) + h(s_k)$ minimal
- 11: **if** $\text{goal}(s_k)$ **then**
- 12: return $s_0 \dots s_k$ (and terminate)
- 13: **else** for every successor s of s_k add $s_0 \dots s_k s$ to frontier
- 14: **end if**
- 15: **end while**

A* Example

Reaching G from S

Recall: $f(s_0 \dots s_k) = g(s_0 \dots s_k) + h(s_k)$

Exp. paths	Frontier
	$\{S : 8\}$




A* Example

Reaching G from S

Recall: $f(s_0 \dots s_k) = g(s_0 \dots s_k) + h(s_k)$

Exp. paths	Frontier
	$\{S : 8\}$
S not goal	$\{SA : 9, SB : 9, SC : 11\}$



Selected path: $S : 8$

Is the last state in S_{goal} ? No


Expand S : add $SA : 1 + 8, SB : 5 + 4, SC : 8 + 3$ to the frontier

A* Example

Reaching G from S

Recall: $f(s_0 \dots s_k) = g(s_0 \dots s_k) + h(s_k)$

Exp. paths	Frontier
	$\{S : 8\}$
S not goal	$\{SA : 9, SB : 9, SC : 11\}$
SA not goal	$\{SB : 9, SC : 11, SAD : \infty, SAE : \infty, SAG : 10\}$



Selected path: $SA : 9$

Is the last state in S_{goal} ? No


Expand SA : add $SAD : 4 + \infty$, $SAE : 8 + \infty$, $SAG : 10 + 0$

A* Example

Reaching G from S

Recall: $f(s_0 \dots s_k) = g(s_0 \dots s_k) + h(s_k)$

Exp. paths	Frontier
	$\{S : 8\}$
S not goal	$\{SA : 9, SB : 9, SC : 11\}$
SA not goal	$\{SB : 9, SC : 11, SAD : \infty, SAE : \infty, SAG : 10\}$
SB not goal	$\{SC : 11, SAD : \infty, SAE : \infty, SAG : 10, SBG : 9\}$



Selected path: $SB : 9$

Is the last state in S_{goal} ? No


Expand SB : add $SBG : 9 + 0$ to the frontier

A* Example

Reaching G from S

Recall: $f(s_0 \dots s_k) = g(s_0 \dots s_k) + h(s_k)$

Exp. paths	Frontier
	$\{S : 8\}$
S not goal	$\{SA : 9, SB : 9, SC : 11\}$
SA not goal	$\{SB : 9, SC : 11, SAD : \infty, SAE : \infty, SAG : 10\}$
SB not goal	$\{SC : 11, SAD : \infty, SAE : \infty, SAG : 10, SBG : 9\}$
SBG goal	$\{SC : 11, SAD : \infty, SAE : \infty, SAG : 10\}$




Selected path: $SBG : 9$

Is the last state in S_{goal} ? Yes!

Path found: SBG with a cost of 9


Romania Example




Straight-line distance
to Bucharest

Arad	366
Bucharest	0
Craiova	160
Dobreta	242
Eforie	161
Fagaras	178
Giurgiu	77
Hirsova	151
Iasi	226
Lugoj	244
Mehadia	241
Neamt	234
Oradea	380
Pitesti	98
Rimnicu Vilcea	193
Sibiu	253
Timisoara	329
Urziceni	80
Vaslui	199
Zerind	374


A* Search Example




A* Search Example




A* Search Example



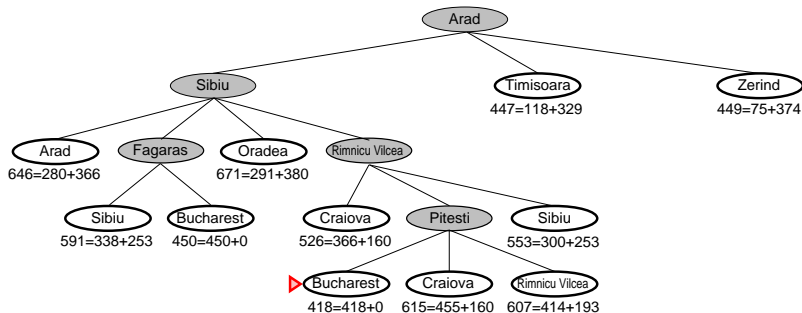
A* Search Example



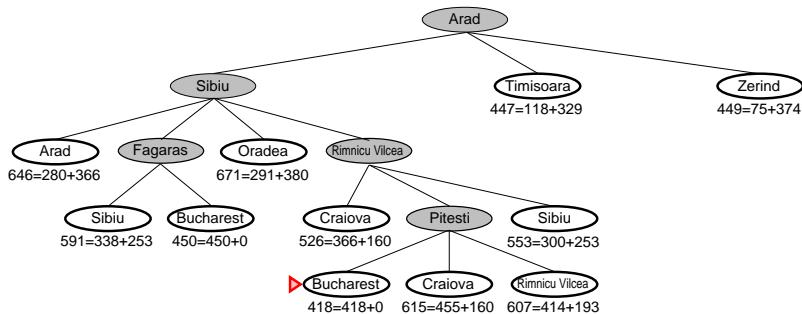
A* Search Example



A* Search Example



A* Search Example



Total distance to go: 418 km!

Properties of A* search

- ▶ Complete and optimal under minor conditions if
 - ▶ an **admissible** heuristic h is used:

$$h(s) \leq h^*(s)$$

where h^* is the true cost from s to a goal.

Properties of A* search

- ▶ Complete and optimal under minor conditions if
 - ▶ an **admissible** heuristic h is used:

$$h(s) \leq h^*(s)$$

where h^* is the true cost from s to a goal.

- ▶ Thus, a heuristic h is admissible if it never overestimates the distance to the goal (is optimistic).

Examples of admissible heuristics for 8-Puzzle

- ▶ $h_1(s)$ = number of misplaced tiles.
- ▶ $h_2(s)$ = Manhattan distance. Take for each tile the sum over the horizontal and vertical steps from the desired location (its Manhattan distance from the desired location). Then take the sum over those distances.

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

$$h_1(s) = ??$$

$$h_2(s) = ??$$

Examples of admissible heuristics for 8-Puzzle

- ▶ $h_1(s)$ = number of misplaced tiles.
- ▶ $h_2(s)$ = Manhattan distance. Take for each tile the sum over the horizontal and vertical steps from the desired location (its Manhattan distance from the desired location). Then take the sum over those distances.

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

$$h_1(s) = 6$$

$$h_2(s) = ??$$

Examples of admissible heuristics for 8-Puzzle

- ▶ $h_1(s)$ = number of misplaced tiles.
- ▶ $h_2(s)$ = Manhattan distance. Take for each tile the sum over the horizontal and vertical steps from the desired location (its Manhattan distance from the desired location). Then take the sum over those distances.

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

$$h_1(s) = 6$$

$$h_2(s) = 4+0+3+3+1+0+2+1 = 14$$

Importance of the Heuristic Choice

Typical search costs (data averaged over 100 instances of the 8-puzzle and d the length of the shortest solution path):

$d = 14$ IDS = 3,473,941 paths

$A^*(h_1) = 539$ paths

$A^*(h_2) = 113$ paths

$d = 24$ IDS \approx 54,000,000,000 paths

$A^*(h_1) = 39,135$ paths

$A^*(h_2) = 1,641$ paths

Summary

- ▶ Heuristic functions estimate costs of shortest paths
- ▶ Good heuristics can dramatically reduce search cost
- ▶ Greedy best-first search expands lowest h
 - ▶ incomplete and not always optimal
- ▶ A* search expands lowest $g + h$
 - ▶ complete and optimal
 - ▶ also optimally efficient